

State-transition model dynamic - Sick-Sicker case study

*Eline Krijkamp**, *MSc*[†]
*Fernando Alarid-Escudero*¹, *PhD*[‡]
Eva A. Enns, *PhD*³
Myriam G.M. Hunink, *PhD*^{2§}
Petros Pechlivanoglou, *PhD*[¶]
Hawre Jalal, *MD*, *PhD*^{||}

2019-02-13

Sick-Sicker model

In the Sick-Sicker model, we simulate a hypothetical cohort of 25-year-old individuals over a lifetime (or reaching age 100-years old) using 75 annual cycles, represented with `n.t`. The cohort start in the “Healthy” health state (denoted H). Healthy individuals are at risk of developing the illness, at which point they would transition to the first stage of the disease (the “Sick” health state, denoted S1). Individuals that become sick incur a one-time utility decrement of 0.01 (`du.HS1`), disutility of transitioning from H to S1) and a one-time cost of \$1,000 (`ic.HS1`) that reflect the acute impacts of developing the illness. Sick individuals are at risk of further progressing to a more severe stage (the Sicker health state, denoted S2), which is constant in this case example. There is a chance that individuals in the Sick state eventually recover and return back to the Healthy state. However, once an individual reaches the Sicker health state, they cannot recover; that is, the probability of transitioning to the “Sick” or “Healthy” health states from the Sicker health state is zero. Individuals in the “Healthy” state face background mortality that is age-specific (i.e., time-dependent). Sick and Sicker individuals face an increased mortality in the form of a hazard rate ratio (HR) of 3 and 10 times, respectively, on the background mortality rate. Sick and Sicker individuals also experience increased health care costs and reduced QoL compared to healthy individuals. Once simulated individuals die, they transition to the Dead health state (denoted D), where they remain. When an individual dies, they incur a one-time cost of \$2,000 (`ic.D`) that reflects the acute care that might be received immediately preceding death. The state-transition diagram of the Sick-Sicker model is shown in Figure 1. The evolution of the cohort is simulated in one-year discrete-time cycles. Both costs and QALYs are discounted at an annual rate of 3%.

01 Initial setup

We start with loading the packages and functions needed.

01.2 External parameters

In the external parameter set up we specify the starting age of the cohort, the number of cycles, the names of the health states and the discount rate used for costs and QALYs. The age specific mortality rate for those in the healthy state is based on the US overall mortality data from the Human Mortality Database. The base-case parameters are combined in a list using the `f.generate_init_params()` function.

```
#### 01.2.1 General setup ####  
n.age.init <- 25 # age of starting cohort
```

*Equal contribution to this paper

[†]Department of Epidemiology, Erasmus MC, University Medical Center, Rotterdam, The Netherlands

[‡]Division of Health Policy and Management, University of Minnesota School of Public Health, Minneapolis, MN, USA

[§]Center for Health Decision Sciences, Harvard T.H. Chan School of Public Health, Boston, USA

[¶]The Hospital for Sick Children, Toronto and University of Toronto, Toronto, Ontario, Canada

^{||}University of Pittsburgh, Pittsburgh, PA, USA




Figure 1: Sick-Sicker

```
n.t      <- 75 # time horizon, number of cycles
v.age.names <- n.age.init:(n.age.init + n.t - 1) # vector with age names
v.n <- c("H", "S1", "S2", "D") # vector with the 4 health states of the model:
# Healthy (H), Sick (S1), Sicker (S2), Dead (D)
n.states <- length(v.n) # number of health states
d.c <- 0.03 # discount rate for costs
d.e <- 0.03 # discount rate for QALYs
v.dwc <- 1 / ((1 + d.e) ^ (0:(n.t))) # vector with discount weights for costs
v.dwe <- 1 / ((1 + d.c) ^ (0:(n.t))) # vector with discount weights for QALYs
v.s.init <- c(H = 1, S1 = 0, S2 = 0, D = 0) # initial state vector

#### 01.2.2 All-cause age-, sex- and race- (ASR) specific mortality ####
df.r.asr <- read.csv("../data/01_all-cause-mortality-USA-2015.csv")
v.r.asr <- df.r.asr %>%
  dplyr::select(Total) %>%
  as.matrix() # vector with mortality rates

#### 01.2.3 Generate initial set of base-case external parameters ####
v.params.init <- f.generate_init_params()
## Create name of parameters
v.names.params <- names(v.params.init)
```

02 Define and initialize matrices and vectors

In this section we initialize the matrices and vectors used for storing the data. The transition probability matrix is initialized and filled for the current cycle using the function `f.create_transition_prob_matrix` with arguments `v.params` and `t`. The function starts with calculating the age-specific transition probabilities for the current cycle `t` and uses these parameters together with the parameters stored in `v.params.init` to fill the matrix. For the initiation of cycle 0 we use `t = 1`. The next step is to initialize the state vector `S0`. In the Sick-Sicker model all individuals start in the Healthy Health state. This state vector is used to inform

the first row of initialized the Markov trace matrix `m.P`.

Equation 1

```
#### 02.1 Transition probability matrix ####
# matrix m.P at the first cycle
f.create_transition_prob_matrix(v.params = v.params.init, t = 1)

##           H           S1           S2           D
## H  0.8489865 0.1500000 0.0000000 0.001013486
## S1 0.5000000 0.3919626 0.1050000 0.003037378
## S2 0.0000000 0.0000000 0.9899112 0.010088764
## D  0.0000000 0.0000000 0.0000000 1.000000000

#### 02.2 Initial state vector ####
# the cohort start in the Healthy health state
s0 <- c(H = 1, S1 = 0, S2 = 0, D = 0)
s0

##   H S1 S2 D
##   1  0  0  0

#### 02.3 Cohort trace
## Create the Markov trace matrix m.M capturing the proportion of the cohort
# in each state at each cycle

m.M <- matrix(0, # initialize cohort trace
              nrow = (n.t + 1), ncol = n.states,
              dimnames = list(0:n.t, v.n))

m.M[1, ] <- s0 # store the initial state vector
```

Now we specified all parameters for the general set up, we specified our input parameters, initialized all structures and fill the transition probability matrix `m.P` and the first row of our Cohort trace `m.M`. This allows us to start running the Markov model.

03 Matrix approach

In this section we show how we can run the Markov model for all cycles. The calculation shown in Equation 2 needs to be performed for all cycles. Therefore, we create a loop starting at `t = 1` until `t = n.t`. Since our transition probabilities are depending on the age of the individuals in the cohort (e.g. we use age dependent mortality rates), we need to update the transition probability matrix `m.P` every cycle. This cycle dependent matrix is used for the matrix multiplication, specified in R with `%*%`, with the cohort trace, `m.M[t,]` to fill the next row of the `m.M[t + 1,]`. ##### Equation 2

```
for(t in 1:n.t){ # loop through the number of cycles
  # create the transition probability matrix for the current cycle
  m.P <- f.create_transition_prob_matrix(v.params = v.params.init, t = t)
  # estimate the state vector for the next cycle (t + 1)
  m.M[t + 1, ] <- m.M[t, ] %*% m.P # Equation 2
}
```

When printing the first six rows of `m.M` we see that everyone starts in the Healthy health state and over time the cohort transitions towards the other three health states. Until cycle 5, the proportion in `S1` is increasing after which it starts decreasing, while the proportion in `S2` and `D` is increasing towards the end of the model. We now ran our model using the Matrix approach and have information about state occupation at each cycle.

This information allows us to apply state rewards (e.g. `c.Healthy`, `u.Healthy`, `u.S1` etc.), but it is not possible to include the transition rewards (e.g. `ic.HS1`, `du.HS1` and `ic.D`). In order to include these rewards we need to know when individuals made the transition. Therefore, in the next section we will explain the Array approach which makes it possible to include these rewards.

```
head(round(m.M, 3)) # show the first six lines of the Markov trace
```

```
##      H    S1    S2    D
## 0 1.000 0.000 0.000 0.000
## 1 0.849 0.150 0.000 0.001
## 2 0.796 0.186 0.016 0.002
## 3 0.769 0.192 0.035 0.004
## 4 0.749 0.191 0.055 0.006
## 5 0.731 0.187 0.074 0.008
```

04 Array approach

The Array approach starts similar as the Matrix approach, meaning that section 01 and 02 are identical for the two approaches. The biggest difference between Matrix and array approach are the dimensions of the structure to store the dynamics of the cohort. While in the Matrix approach, we stored all information in matrix `m.M` of size `n.states` x `n.states`, in this approach we add an dimension for time resulting in an array with dimensions `n.states` x `n.states` x `n.t`. In R indexing start at 1, therefore, we initialize the array `a.A` using `n.state + 1` to allow storing the results from cycle 0 until cycle `n.t`. The initial state vectors `s0` is used to inform the initial cycle of the array.

```
a.A <- array(0, dim = c(n.states, n.states, n.t + 1),
             dimnames = list(v.n, v.n, 0:n.t)) # initialize array
diag(a.A[, , 1]) <- s0 # store the initial state vector in the diagonal of A
```

Equation 3 & 4

We can now run the model using the Array approach. The function `f.create_transition_prob_matrix` is the same for both the Matrix and Array approach and needs to be calculated for each cycle. This `m.P` for the current cycle is multiplied with array `a.A` using element-wise multiplication `*`. The information about all transitions dynamics is stored in the next cycle in `a.A`.

```
a.A[, , 1]

##      H S1 S2 D
## H    1  0  0  0
## S1    0  0  0  0
## S2    0  0  0  0
## D    0  0  0  0

# run the model
for(t in 1:n.t){
  # loop through the number of cycles
  # create the transition probability matrix for the current cycle
  m.P <- f.create_transition_prob_matrix(v.params = v.params.init, t = t)
  ##### Equation 4 #####
  a.A[, , t + 1] <- colSums(a.A[, , t]) * m.P # fill array A for t + 1
}
```

Equation 5

To get an idea about how the information in `a.A` looks like we print it first three cycles. Like in the transition probability matrix `m.P`, the rows specify in which health state the individual started at the beginning of the cycle, while the columns inform you about where individual transitioned to. In cycle 0 everyone started in the healthy health states. At cycle 1 we can see that 0.849 of the cohort stayed healthy, 0.15 transitioned from Healthy to Sick and `r round(a.A["H","D", "1"], 4)` died. Indeed, this looks very similar to the transition probabilities in this case example. From cycle 2 and onwards the information in `a.A` becomes more interesting. In cycle 2, we see that 0.7208 of the cohort stayed healthy, 0.1273 transitioned from Healthy towards Sick and 0.00084 died while they started out healthy. In addition we see that 0.075 of the cohort recovered from Healthy. 0.0588 stayed Sick, 0.0158 became Sicker out of Sick and 0.00045 died from Sick. All these values sum to 1 since we are still describing what happens to the cohort over time.

```
a.A[, , 1:3] # shown for two cycles
```

```
## , , 0
##
##      H S1 S2 D
## H   1  0  0  0
## S1  0  0  0  0
## S2  0  0  0  0
## D   0  0  0  0
##
## , , 1
##
##           H      S1 S2           D
## H  0.8489865 0.15  0 0.001013486
## S1 0.0000000 0.00  0 0.000000000
## S2 0.0000000 0.00  0 0.000000000
## D  0.0000000 0.00  0 0.000000000
##
## , , 2
##
##           H           S1      S2           D
## H  0.7207908 0.12734798 0.00000 0.000847714
## S1 0.0750000 0.05880112 0.01575 0.000448877
## S2 0.0000000 0.00000000 0.00000 0.000000000
## D  0.0000000 0.00000000 0.00000 0.001013486
```

```
sum(a.A[, , 3])
```

```
## [1] 1
```

Equation 7

When you sum the values in a column of `a.A`, e.g. Sick, you get which proportion of the cohort was in Sick at the end of that cycle.

```
sum(a.A[, "S1", 3])
```

```
## [1] 0.1861491
```

By using the `colSums` function, summing over all columns of `a.A` we can do this for all points in time and when we transpose these results we get the traditional cohort trace `m.M`. Here names `m.M_A` to indicate it is generated via the Array approach.

```
# calculating M from A
m.M_A <- t(colSums(a.A)) # sum over the columns of a.A and transpose
```

Since a Markov model is stochastic, these two approaches should give identical results. We check this using the `==` function. We use rounding on 10 decimals, to avoid wrong `FALSE` results that have to do with floating point comparison issues. This means that functions allow you to test for equality with a specified difference tolerance.

05 Apply state and transition rewards

```
#### 05.1 Create reward matrices for both costs and effects ####
m.R_costs    <- f.create_transition_reward_matrix_costs(v.params = v.params.init)
m.R_effects  <- f.create_transition_reward_matrix_effects(v.params = v.params.init)
```

Equation 8

```
m.R_costs    # show the reward matrix for costs

##          H      S1      S2      D
## H      2000  3000  2000  4000
## S1     4000  4000  4000  6000
## S2    15000 15000 15000 17000
## D           0      0      0      0

m.R_effects  # show the reward matrix for effects

##          H      S1      S2      D
## H      1.00  0.99  1.00  1.00
## S1     0.75  0.75  0.75  0.75
## S2     0.50  0.50  0.50  0.50
## D      0.00  0.00  0.00  0.00
```

Equation 9

```
#### 05.2 Expected QALYs and Costs per cycle for each strategy ####
a.0_costs <- a.0_effects <- array(0, dim = c(n.states, n.states, n.t + 1),
                                   dimnames = list(v.n, v.n, 0:n.t))

for(t in 1:n.t){
  # element-wise-multiplication of array A with the rewards matrices
  a.0_costs[, , t] <- a.A[, , t] * m.R_costs
  a.0_effects[, , t] <- a.A[, , t] * m.R_effects
}
```

Equation 10

The final step is to calculate the total expected discounted costs and QALYs. We start by calculating the expected cost and QALYs per cycle. These values, stored in the vectors `v.Costs` and `v.QALYs`, in turn are multiplied with the vector of discount weights, `v.dwc1` and `v.dwe`, respectively. This gives us the total expected discounted cost (TC) and QALYs (TE).

```
# calculate the expected costs per cycle
v.Costs <- rowSums(t(colSums(a.0_costs)))
# calculate the expected QALYs per cycle
v.QALYs <- rowSums(t(colSums(a.0_effects)))

TC <- t(v.Costs) %*% v.dwc # calculate the total expected discounted costs
```

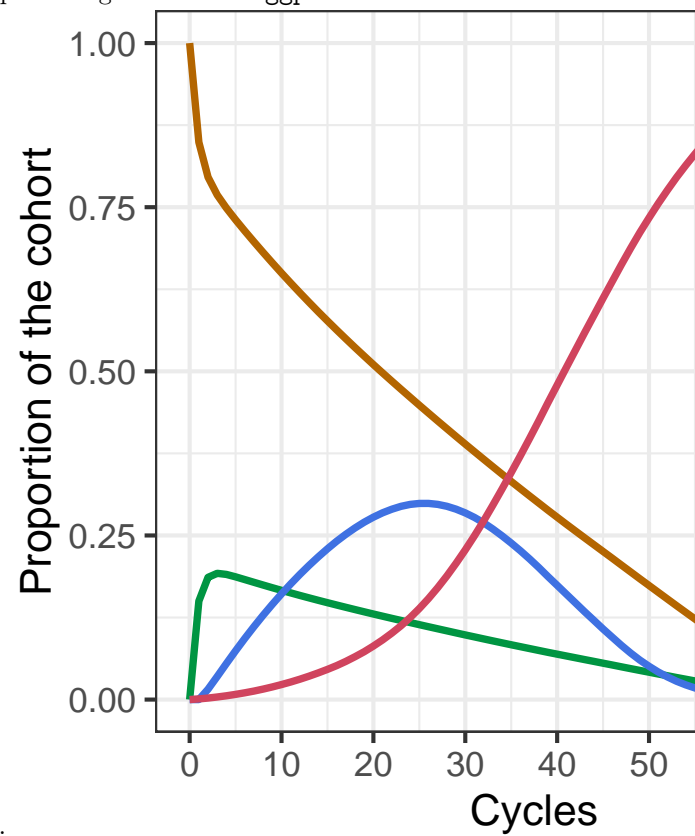
```
TE <- t(v.QALYs) %*% v.dwe # calculate the total expected discounted QALYs

v.Results <- c(TC, TE)      # combine the total expected costs and QALYs
names(v.Results) <- c("Costs", "Effect") # name the vector
v.Results                  # print the results
```

```
##          Costs          Effect
## 115104.13376      20.37685
```

06 Plot cohort trace

The results of a cohort trace are much easier to interpret via a graph. Using the function `ggplot` can show the



proportion of the cohort in each state (y-axis) at each cycle (x-axis).